

MCS 7103_Assignment_one

September 11, 2024

NAME: JOSHUA Byaruhanga REG. NO.: 2024/HD05/21925U STUDENT NO.: 2400721925

Predicting if a client will subscribe to a Term Deposit

Business problem:

There has been a revenue decline for the Portuguese bank, and they would like to know what actions to take. After investigation, we found out that the root cause is that their clients are not depositing as frequently as before. Knowing that term deposits allow banks to hold onto a deposit for a specific amount of time, so banks can invest in higher gain financial products to make a profit. In addition, banks also hold a better chance to persuade term deposit clients into buying other products such as funds or insurance to further increase their revenues. As a result, the Portuguese bank would like to identify existing clients that have a higher chance to subscribe for a term deposit and focus marketing effort on such clients.

Data Wrangling

The Data to be used to predict if a client will subscribe to A Term Deposit was obtained from <https://archive.ics.uci.edu/dataset/222/bank+marketing>

A classification approach is used to predict which clients are more likely to subscribe for term deposits.

```
[24]: # import required python libray
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
[25]: #Read data set
df=pd.read_csv("bank-full.csv",delimiter=';',quotechar='')
```

Printing the first 5 rows of the dataset

```
[26]: df.head()
```

```
[26]:   age      job  marital  education  default  balance  housing  loan  \
0   58  management  married   tertiary     no     2143     yes    no
1   44  technician  single   secondary     no      29     yes    no
2   33  entrepreneur  married   secondary     no      2     yes    yes
```

3	47	blue-collar	married	unknown	no	1506	yes	no
4	33	unknown	single	unknown	no	1	no	no

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no

What's the size of Dataset

```
[27]: df.shape
```

```
[27]: (45211, 17)
```

The dataset consist of 45211 rows and 17 columns

What are the different fields and respective Datatypes that constitute the Dataset

```
[28]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         45211 non-null  int64
1   job         45211 non-null  object
2   marital     45211 non-null  object
3   education   45211 non-null  object
4   default     45211 non-null  object
5   balance     45211 non-null  int64
6   housing     45211 non-null  object
7   loan        45211 non-null  object
8   contact     45211 non-null  object
9   day         45211 non-null  int64
10  month       45211 non-null  object
11  duration    45211 non-null  int64
12  campaign    45211 non-null  int64
13  pdays       45211 non-null  int64
14  previous    45211 non-null  int64
15  poutcome    45211 non-null  object
16  y           45211 non-null  object
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

Age: This is age of client of Numeric type

Job : This is client's title of categorical type

Marital: This is marital status for the customer(Categorical)

Education: This field shows the education level for the client(Categorical)

default: whether customer has defaulted credit or not(Binary)

balance: customer average yearly balance(Numeric) housing: whether a customer has a housing loan or not (Binary)

loan: Whether customer has personal loan or not (Binary)

contact: whether customer was contacted using cellular or telephone(categorical)

day: Last day of the month (Numeric)

Month: last contact month of year (categorical)

duration: Last contact duration, in seconds (numeric)

campaign: number of contacts performed during this campaign and for this client (numeric)

pdays: number of days that passed by after the client was last contacted from a previous campaign (numeric, -1 means client was not previously contacted)

previous: number of contacts performed before this campaign and for this client (numeric)

poutcome: outcome of the previous marketing campaign (categorical: “unknown”, “other”, “failure”, “success”)

y: has the client subscribed a term deposit? (binary: “yes”, “no”)

find out duplication

```
[29]: df.nunique()
```

```
[29]: age          77
      job          12
      marital      3
      education    4
      default      2
      balance     7168
      housing      2
      loan         2
      contact      3
      day         31
      month       12
      duration    1573
      campaign     48
      pdays       559
      previous     41
      poutcome     4
      y           2
      dtype: int64
```

Find out missing values in the dataset

```
[30]: df.isnull().sum()
```

```
[30]: age          0
      job          0
      marital     0
      education   0
      default     0
      balance     0
      housing     0
      loan        0
      contact     0
      day         0
      month       0
      duration    0
      campaign    0
      pdays       0
      previous    0
      poutcome    0
      y          0
      dtype: int64
```

EDA Exploratory Data Analysis

Univariate Analysis

Analysis of 'Subscribed' variable

```
[32]: #Frequency of 'subscribed'
      df['y'].value_counts()
```

```
[32]: y
      no      39922
      yes     5289
      Name: count, dtype: int64
```

```
[33]: df1=df
```

```
[34]: df1
```

```
[34]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...	
45206	51	technician	married	tertiary	no	825	no	no	
45207	71	retired	divorced	primary	no	1729	no	no	

45208	72	retired	married	secondary	no	5715	no	no
45209	57	blue-collar	married	secondary	no	668	no	no
45210	37	entrepreneur	married	secondary	no	2971	no	no

	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	no
1	unknown	5	may	151	1	-1	0	unknown	no
2	unknown	5	may	76	1	-1	0	unknown	no
3	unknown	5	may	92	1	-1	0	unknown	no
4	unknown	5	may	198	1	-1	0	unknown	no
...
45206	cellular	17	nov	977	3	-1	0	unknown	yes
45207	cellular	17	nov	456	2	-1	0	unknown	yes
45208	cellular	17	nov	1127	5	184	3	success	yes
45209	telephone	17	nov	508	4	-1	0	unknown	no
45210	cellular	17	nov	361	2	188	11	other	no

[45211 rows x 17 columns]

Converting the target variables into 0s and 1s

```
[35]: df1['y'] = df1['y'].map({'yes': 1, 'no': 0})
```

```
[36]: df1
```

```
[36]:
```

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	unknown	no	1506	yes	no	
4	33	unknown	single	unknown	no	1	no	no	
...
45206	51	technician	married	tertiary	no	825	no	no	
45207	71	retired	divorced	primary	no	1729	no	no	
45208	72	retired	married	secondary	no	5715	no	no	
45209	57	blue-collar	married	secondary	no	668	no	no	
45210	37	entrepreneur	married	secondary	no	2971	no	no	

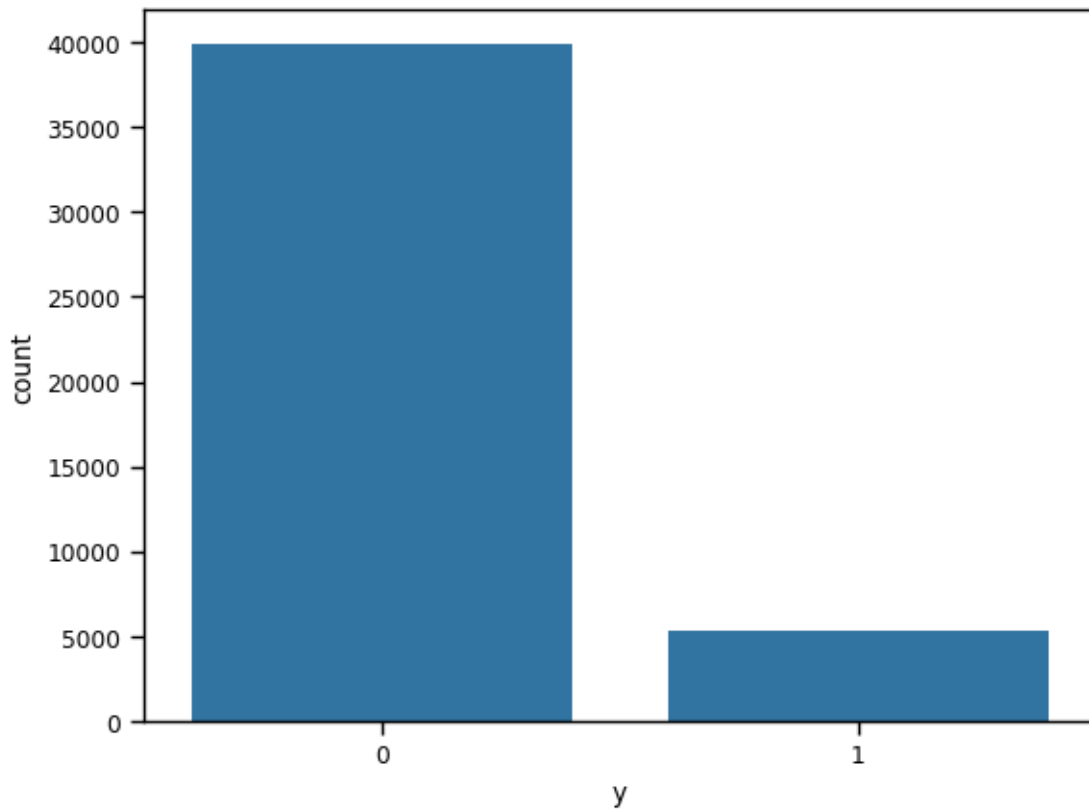
	contact	day	month	duration	campaign	pdays	previous	poutcome	y
0	unknown	5	may	261	1	-1	0	unknown	0
1	unknown	5	may	151	1	-1	0	unknown	0
2	unknown	5	may	76	1	-1	0	unknown	0
3	unknown	5	may	92	1	-1	0	unknown	0
4	unknown	5	may	198	1	-1	0	unknown	0
...
45206	cellular	17	nov	977	3	-1	0	unknown	1
45207	cellular	17	nov	456	2	-1	0	unknown	1

45208	cellular	17	nov	1127	5	184	3	success	1
45209	telephone	17	nov	508	4	-1	0	unknown	0
45210	cellular	17	nov	361	2	188	11	other	0

[45211 rows x 17 columns]

```
[46]: # Plotting the 'subscribed' frequency
sns.countplot(data=df1, x=df1['y'])
```

```
[46]: <Axes: xlabel='y', ylabel='count'>
```



```
[47]: df1['y'].value_counts(normalize=True)
```

```
[47]: y
0    0.883015
1    0.116985
Name: proportion, dtype: float64
```

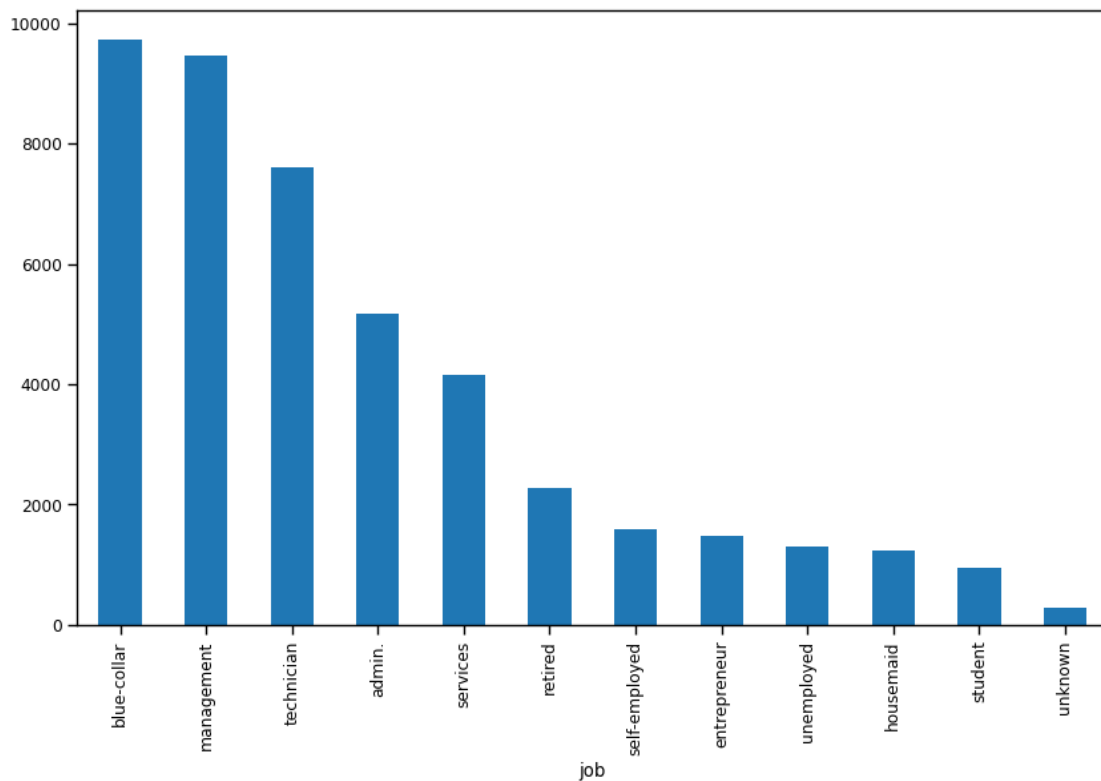
From the above analysis we can see that roughly 12% subscribed to term deposit

Analysing th 'Job' variable

```
[42]: #Frequency table  
df1['job'].value_counts()
```

```
[42]: job  
blue-collar      9732  
management      9458  
technician       7597  
admin.           5171  
services         4154  
retired          2264  
self-employed    1579  
entrepreneur     1487  
unemployed       1303  
housemaid        1240  
student          938  
unknown          288  
Name: count, dtype: int64
```

```
[43]: # Plotting the job frequency table  
sns.set_context('paper')  
df1['job'].value_counts().plot(kind='bar', figsize=(10,6));
```




Most of the clients belonged to blue-collar job and students are least in general as they don't make term deposits in general.

Analysis of 'marital' status

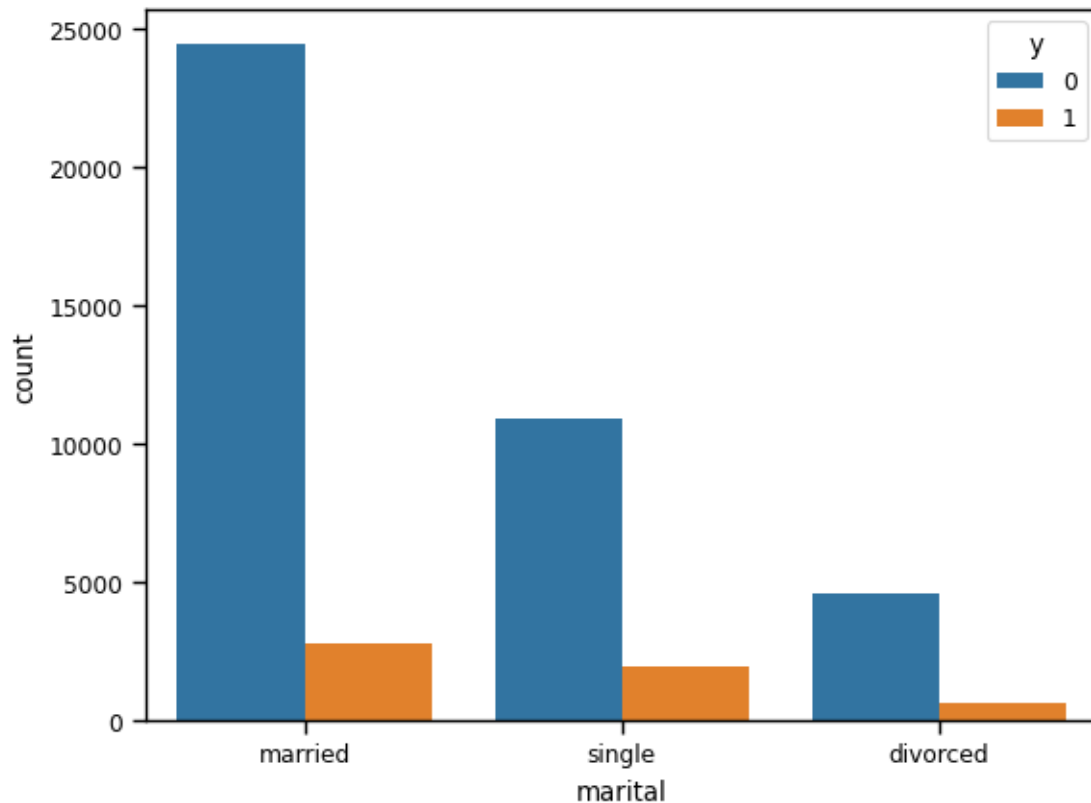
```
[44]: df1['marital'].value_counts()
```

```
[44]: marital  
married    27214  
single     12790  
divorced    5207  
Name: count, dtype: int64
```

```
[45]: sns.countplot(data=df1, x='marital');
```



```
[49]: sns.countplot(data=df1, x='marital', hue='y');
```

from the above graph, married subscribed more to term deposit, followed by single, lastly divorced
Analyzing the 'age' variable

```
[50]: sns.distplot(df1['age']);
```

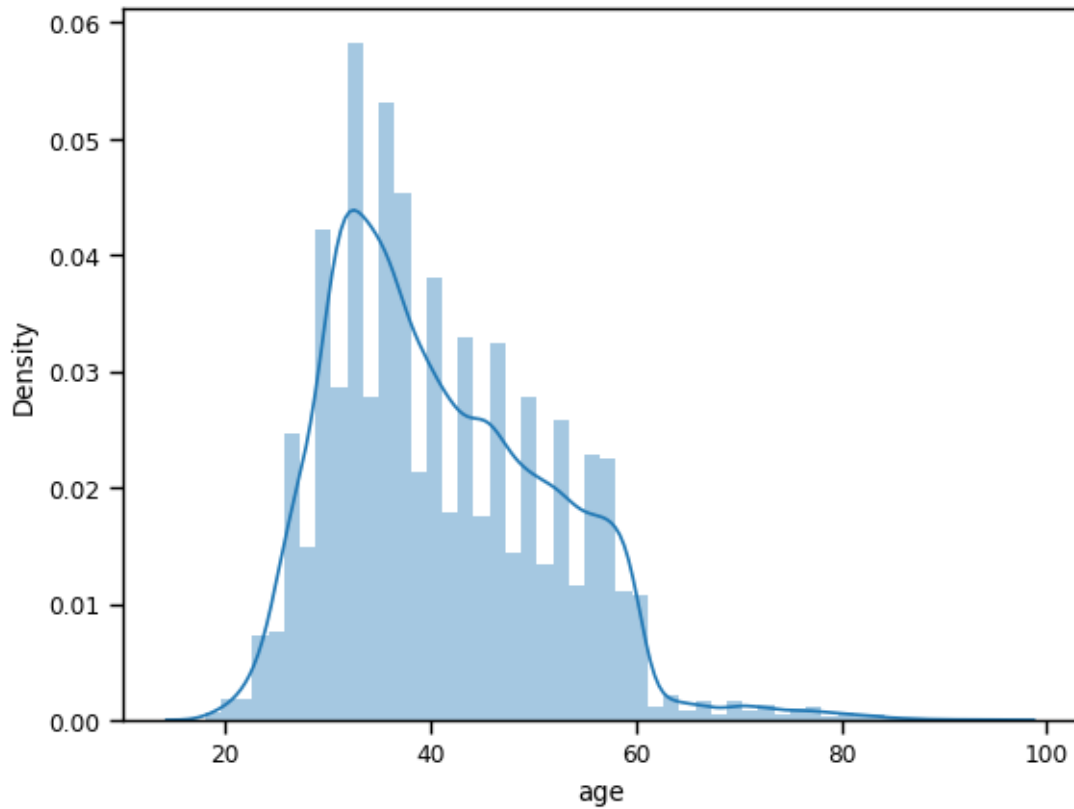
/tmp/ipykernel_5136/1708864848.py:1: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

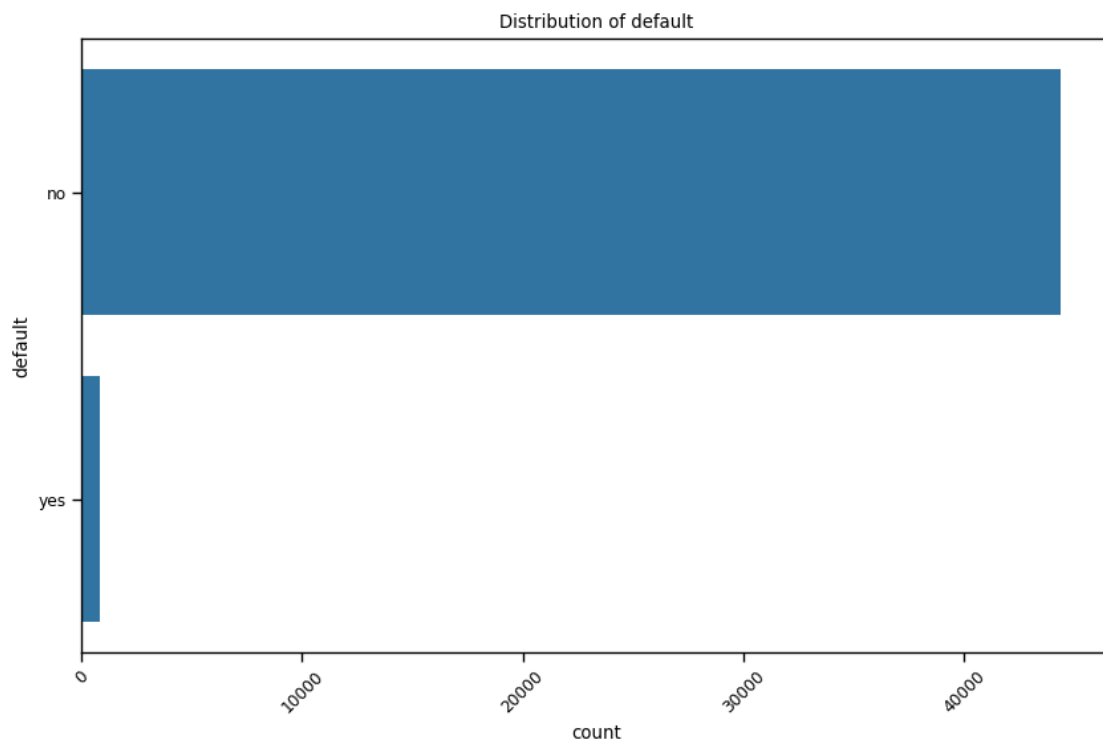
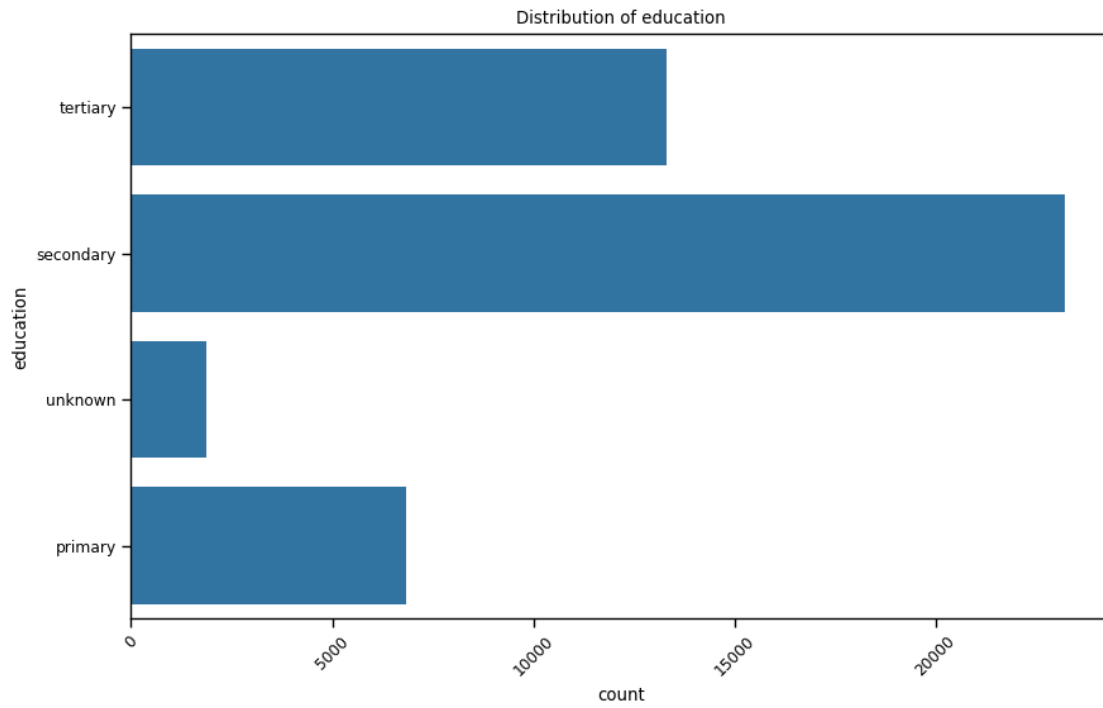
```
sns.distplot(df1['age']);
```

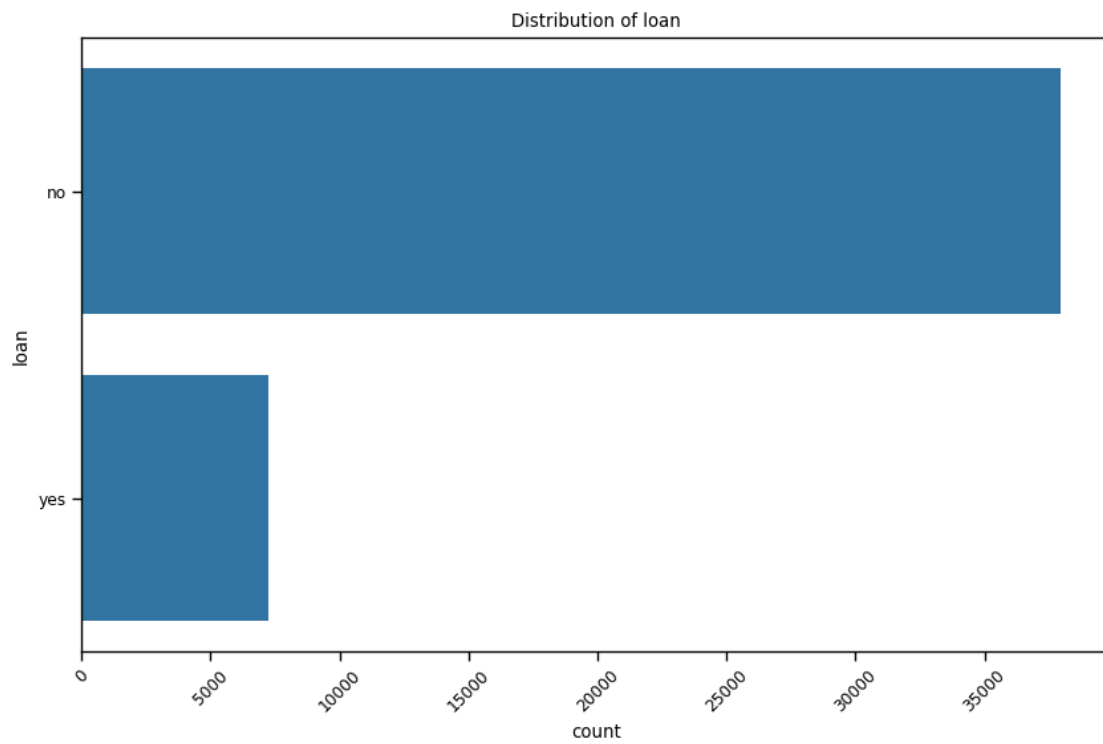
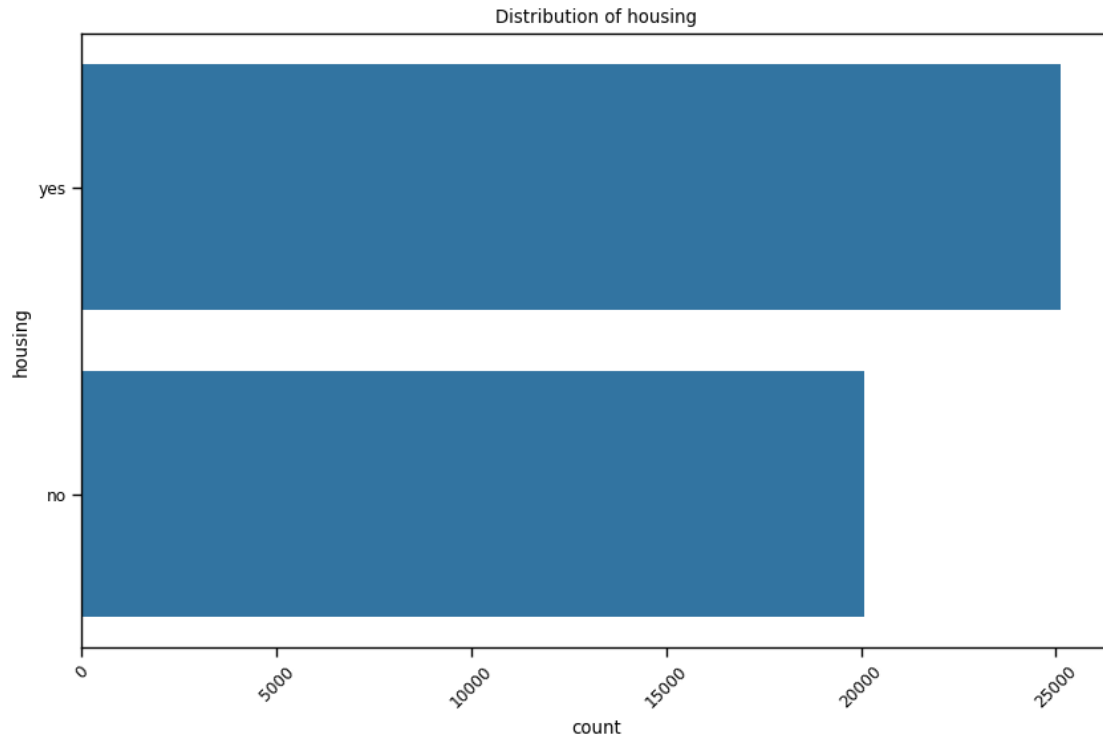


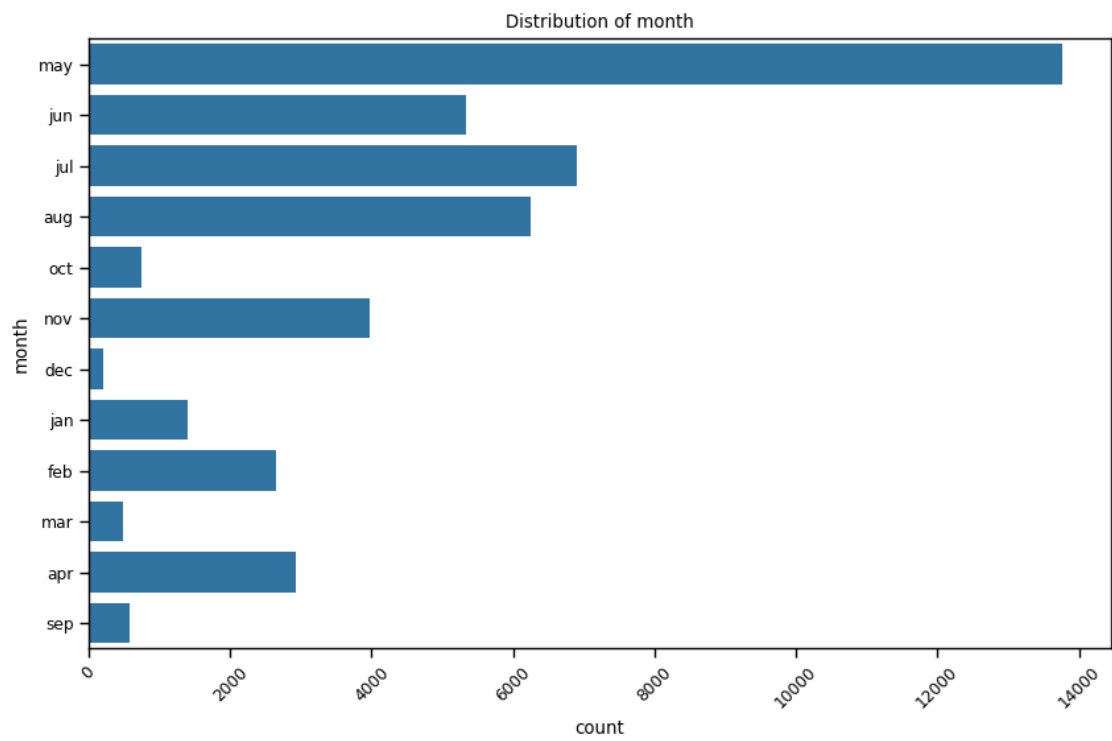
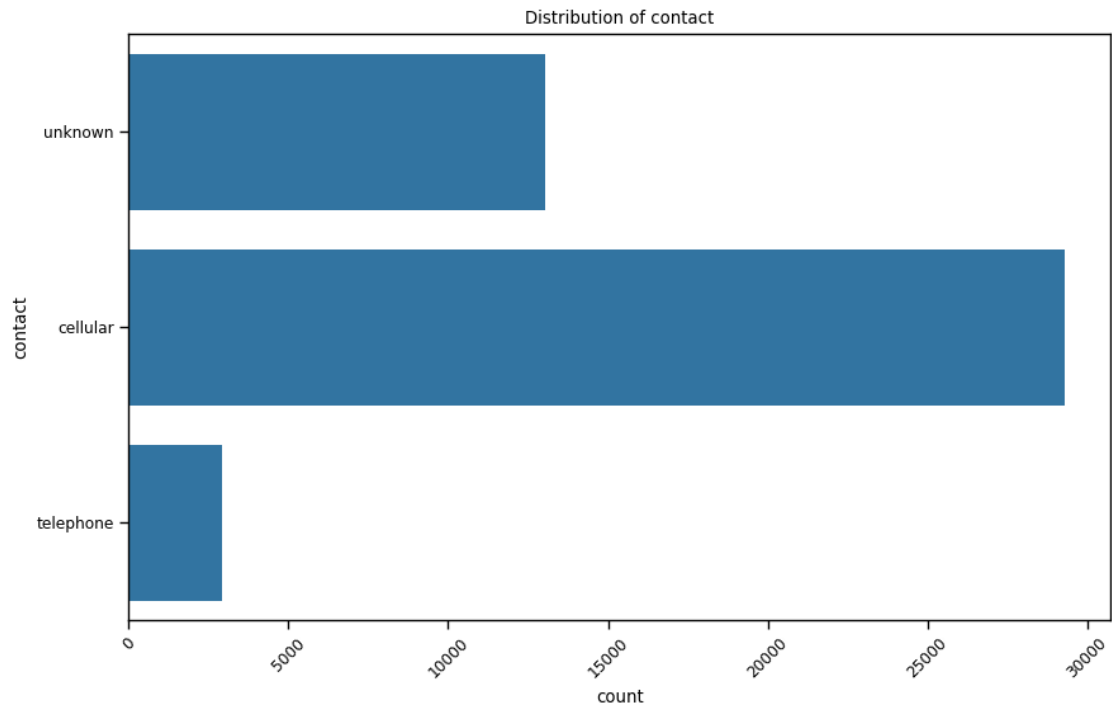
Most of the clients fall in the age group between 20-60.

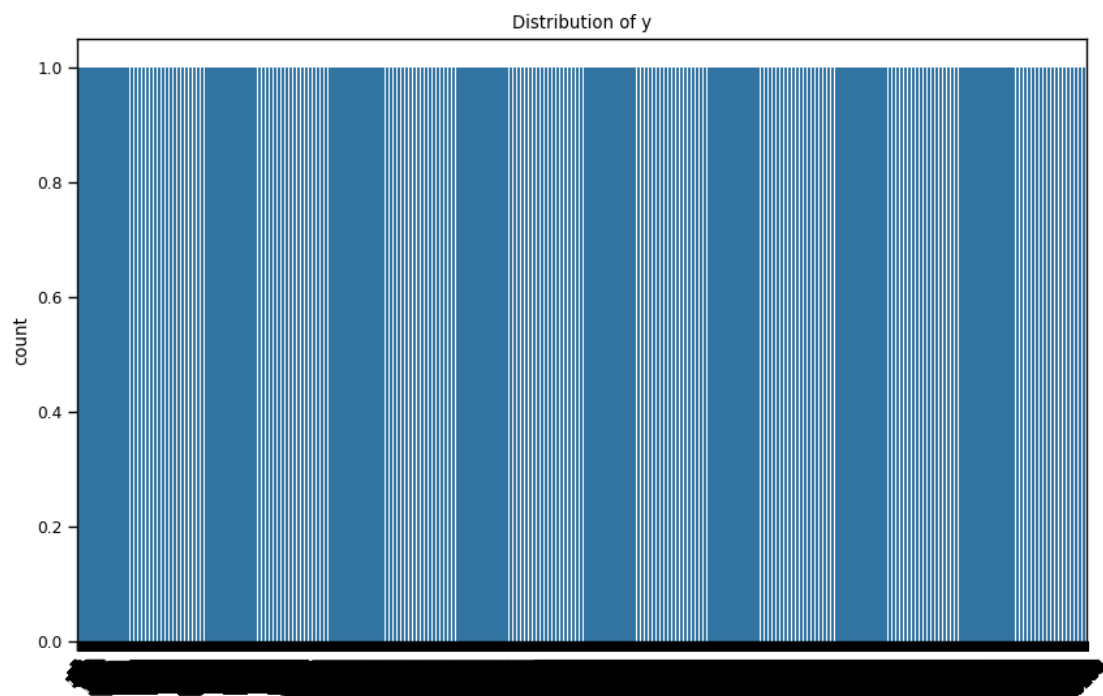
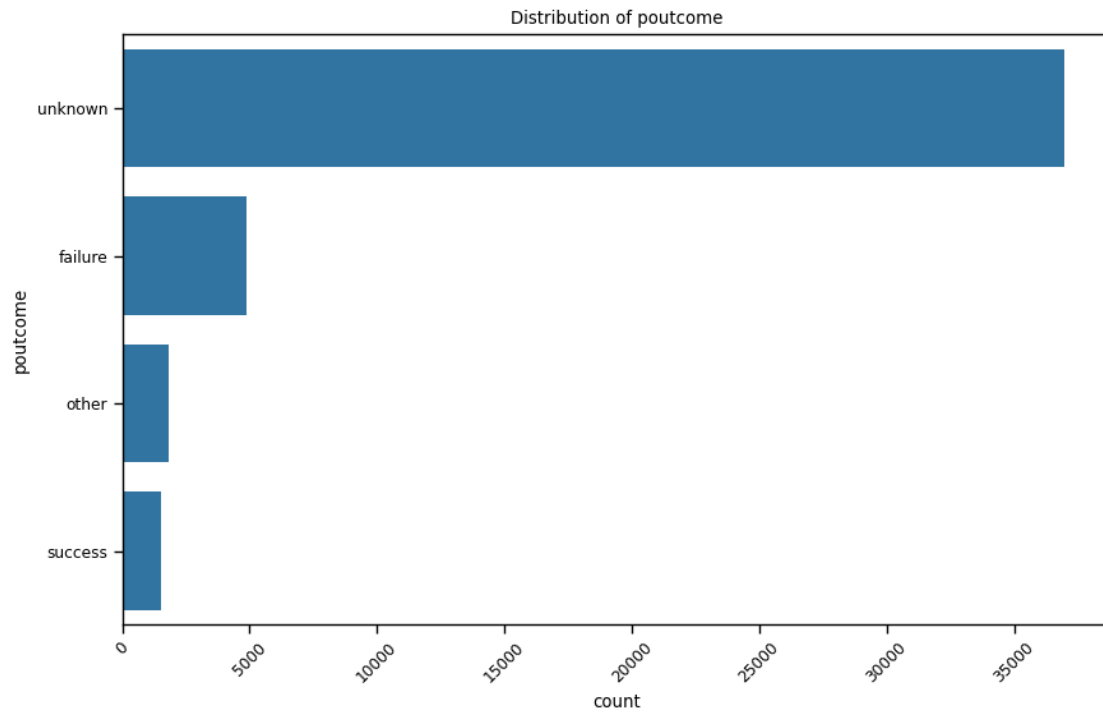
Distribution of other categorical values

```
[74]: categorical_features = [ 'education', 'default', 'housing', 'loan', 'contact', '
    ↪ 'month', 'poutcome', 'y']
for feature in categorical_features:
    plt.figure(figsize=(10, 6))
    sns.countplot(df[feature])
    plt.title(f'Distribution of {feature}')
    plt.xticks(rotation=45)
    plt.show()
```










Bivariate Analysis

```
[52]: #job vs subscribed
print(pd.crosstab(df1['job'],df1['y']))
```

y	0	1
job		
admin.	4540	631
blue-collar	9024	708
entrepreneur	1364	123
housemaid	1131	109
management	8157	1301
retired	1748	516
self-employed	1392	187
services	3785	369
student	669	269
technician	6757	840
unemployed	1101	202
unknown	254	34

```
[54]: job = pd.crosstab(df1['job'],df1['y'])
job_norm = job.div(job.sum(1).astype(float), axis=0)
```

```
[55]: job_norm.plot.bar(stacked=True,figsize=(8,6));
```



students and retired people have higher chances of subscribing to a term deposit, which is surprising as students generally do not subscribe to a term deposit. The possible reason is that the number of students in the dataset is less and comparatively to other job types, more students have subscribed to a term deposit.

```
[56]: #Marital status vs subscribed
pd.crosstab(df1['marital'], df1['y'])
```

```
[56]: y          0      1
marital
divorced   4585   622
married   24459  2755
single    10878  1912
```

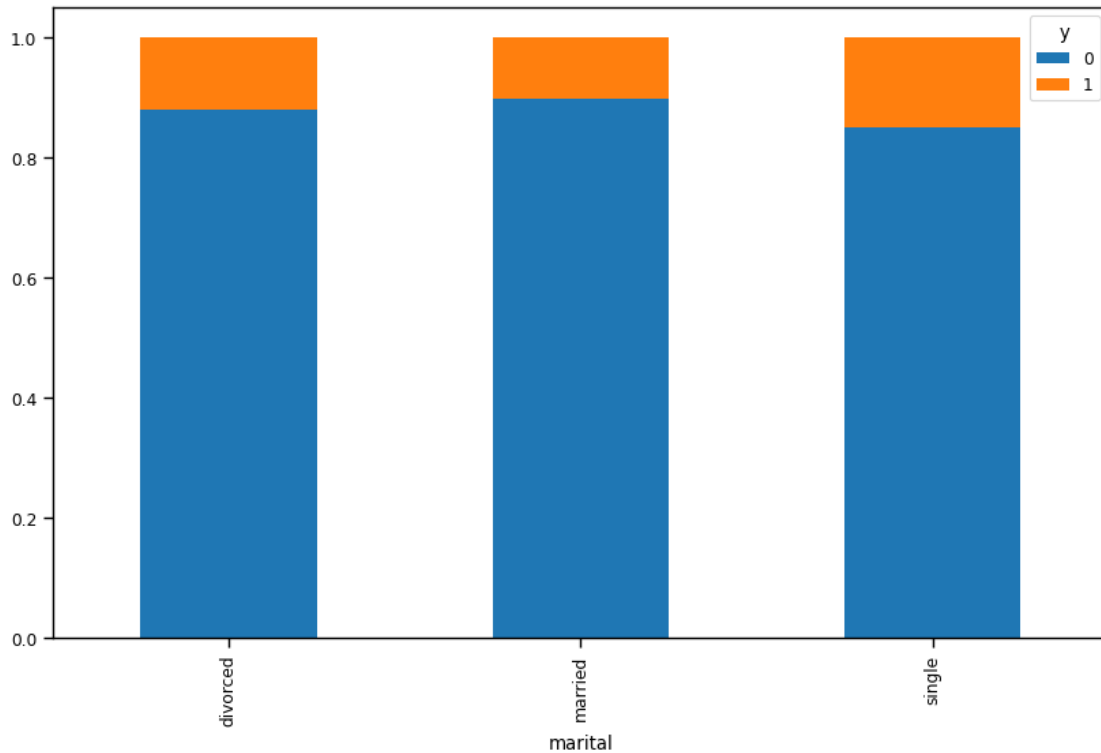
```
[57]: marital = pd.crosstab(df['marital'], df['y'])
marital_norm = marital.div(marital.sum(1).astype(float), axis=0)
```



```
marital_norm
```

```
[57]: y          0          1
      marital
      divorced  0.880545  0.119455
      married   0.898765  0.101235
      single    0.850508  0.149492
```

```
[58]: marital_norm.plot.bar(stacked=True, figsize=(10,6));
```



From the above analysis we can infer that marital status doesn't have a major impact on the subscription to term deposits.

```
[60]: #default vs subscription
      pd.crosstab(df1['default'], df1['y'])
```

```
[60]: y          0          1
      default
      no       39159  5237
      yes        763    52
```


```
[62]: dflt = pd.crosstab(df1['default'], df1['y'])
      dflt_norm = dflt.div(dflt.sum(1).astype(float), axis=0)
```

```
dflt_norm
```

```
[62]: y          0          1  
      default  
no      0.882039  0.117961  
yes      0.936196  0.063804
```

```
[63]: dflt_norm.plot.bar(stacked=True, figsize=(6,6))
```

```
[63]: <Axes: xlabel='default'>
```



clients having no previous default have slightly higher chances of subscribing to a term loan as compared to the clients who have previous default history.

```
[67]: #Correlation matrix  
      corr_column=['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous', 'y']
```

```
cor = df1[corr_column].corr()
cor
```

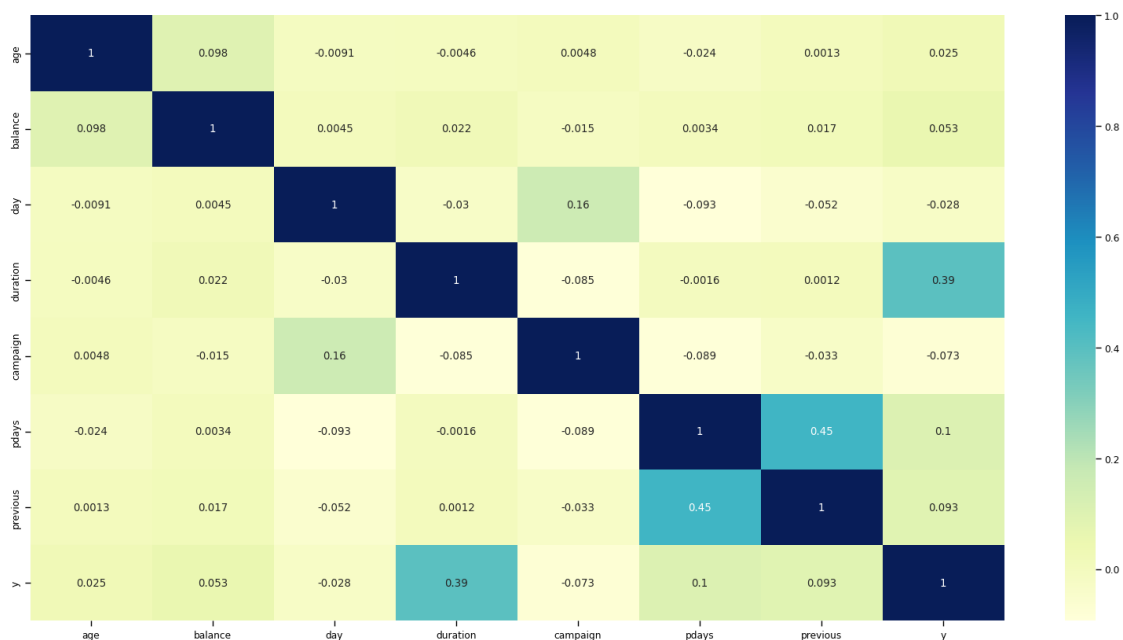
```
[67]:
```

	age	balance	day	duration	campaign	pdays	\
age	1.000000	0.097783	-0.009120	-0.004648	0.004760	-0.023758	
balance	0.097783	1.000000	0.004503	0.021560	-0.014578	0.003435	
day	-0.009120	0.004503	1.000000	-0.030206	0.162490	-0.093044	
duration	-0.004648	0.021560	-0.030206	1.000000	-0.084570	-0.001565	
campaign	0.004760	-0.014578	0.162490	-0.084570	1.000000	-0.088628	
pdays	-0.023758	0.003435	-0.093044	-0.001565	-0.088628	1.000000	
previous	0.001288	0.016674	-0.051710	0.001203	-0.032855	0.454820	
y	0.025155	0.052838	-0.028348	0.394521	-0.073172	0.103621	

	previous	y
age	0.001288	0.025155
balance	0.016674	0.052838
day	-0.051710	-0.028348
duration	0.001203	0.394521
campaign	-0.032855	-0.073172
pdays	0.454820	0.103621
previous	1.000000	0.093236
y	0.093236	1.000000

```
[68]: fig,ax= plt.subplots()
fig.set_size_inches(20,10)
sns.heatmap(cor, annot=True, cmap='YlGnBu')
```

```
[68]: <Axes: >
```



The duration of the call is highly correlated with the target variable. As the duration of the call is more, there are higher chances that the client is showing interest in the term deposit and hence there are higher chances that the client will subscribe to term deposit.

Conclution:

More Job types are Admin, Technician, and blue-collar and it means bank targeting high salaried people.

Most customer were contacted using cellular

Most of the clients fall in the age group between 20-60.

12% of the total client subscribed to term deposit

most of the clients belonged to blue-collar job and students are least in general as they don't make term deposits in general.

clients having no previous default have slightly higher chances of subscribing to a term loan as compared to the clients who have previous default history.

students and retired people have higher chances of subscribing to a term deposit, which is surprising as students generally do not subscribe to a term deposit. The possible reason is that the number of students in the dataset is less and comparatively to other job types, more students have subscribed to a term deposit.

[]: